

Tutorial: x86_64 Assembly Language Programming

by Jerry McIntosh

INTRODUCTION

In this tutorial I hope to share some insight into Assembly language programming by working through developing a list library. That is, a glorified array with some structured bells and whistles. I also hope that this tutorial will be useful as a guide for how to build a shared library in Assembly and C. There will definitely be some insight given into how the two languages work together. And maybe some hint as to the nature of Assembly language. I learn best by doing. Therefore, my hope is that by working through developing a shared library you might gain some insight into x86_64 Assembly language, function calls, parameter passing, and return values. There's some basic math involved to compute offsets and addresses. You'll write some simple iterator functions and see them implemented in test programs written in C. All of the functions in the shared library are written in Assembly language, but have corresponding C declarations that make it possible for the C programs to use them. I provide the output generated by each test program. But, as you progress through this tutorial I suggest you look at the source code of the test programs, and modify them to learn more.

First, here is a list of the requirements for this project.

- Linux OS
- Programming languages: C and Assembly
- Netwide Assembler (NASM) and the GCC compiler
- your favorite text editor
- and working at the command line

We will be writing a shared library in Assembly and C. We will use C to demo our project as we add the bells and whistles that make the list library what it is. Using C will also provide the opportunity to see how the two languages interact.

Assembly language relies heavily on registers. For this project we need access to 64-bit and 32-bit registers. However, there are a number of registers at our disposal: general-purpose registers; the eflags register; and the instruction pointer register.

The 64-bit general-purpose registers are: RAX, RBX, RCX, RDX, RBP, RSP, RSI, RDI, and R8 thru R15.

The 32-bit general-purpose registers are: EAX, EBX, ECX, EDX, EBP, ESP, ESI, EDI, and R8D thru R15D. These registers are located in the lower half (32-bits) of the corresponding 64-bit register. For instance, EAX is the lower 32-bits of RAX, and R8D is the lower 32-bits of R8.

There are also 16-bit and 8-bit registers that form the lower half of the 32-bit and 16-bit registers accordingly. For more information on the general-purpose registers see section "BASIC PROGRAM EXECUTION REGISTERS" in chapter 3 of the "Intel 64 and IA-32 Architectures Software Developer's Manual."

On Linux systems the first 6 parameters (1-6) passed to a function are passed (as listed) in registers: RDI, RSI, RDX, RCX, R8, R9. The remaining parameters (7 and above) are passed on the stack in reverse order. That is the last parameter is pushed on the stack first and the 7th parameter is pushed last.

On Linux systems, functions have a responsibility to preserve the contents of certain registers referred to as callee-saved registers. Here are the callee-saved registers: RBX, RBP, RSP, R12, R13, R14, R15.

We will use the ELF (Executable and Linkable Format) format when generating the object files. The ELF format is used by Linux. You will create the makefile for the shared library, but all other makefiles will be provided. Other than providing the commands necessary to execute the makefiles, I won't be covering the make utility.

NOTE: The Assembly syntax we will use is the Intel syntax.

THE FILE: UTIL.ZIP

If you haven't already downloaded file `util.zip`, do so now. Once downloaded copy the file into your project folder. Next, at the command line change directory to your project folder and unzip the file. Once unzipped the following folder structure should appear.

```
util/  
+ list_add/  
+ list_begin/  
+ list_delete/  
+ list_find/  
+ list_init/  
+ list_sort/  
+ util/
```

NOTE: In the `util/util/` folder you will see file `.gitignore`. As the name suggests, you can ignore the file. GitHub tracks files not folders, so, an empty folder is automatically scrubbed from a repository. You can either ignore the file or delete the file with the following command (while in the `util/util/` folder).

```
rm -f ../.gitignore
```

FUNCTION: MEMMOVE64

In the `util/util/` folder create the file `memmove64.asm` via your text editor, or at the command line with the following command.

```
touch memmove64.asm
```

Now fire up your text editor, open file `memmove64.asm`, and enter the code below.

```
; file:    memmove64.asm  
; brief:   move quadword (8-byte) chunks from source to destination
```

```

%ifndef MEMMOVE64_ASM
#define MEMMOVE64_ASM 1
;
QW_SIZE EQU 8
;
%endif

```

MEMMOVE64_ASM is a marco that is undefined when the NASM assembler first processes file `memmove64.asm`. The NASM directive `%ifndef` (which means: if not defined) controls the processing of the code between the directives `%ifndef` and `%endif`. These directives create what is know as a conditional-block of code. If the marco MEMMOVE64_ASM is undefined then the NASM assembler will process the code between `%ifndef` and the corresponding `%endif`. Within the conditional-block is a `%define` directive that defines macro MEMMOVE64_ASM. Any attempt by the assembler to process the conditional-block after MEMMOVE64_ASM is defined will fail. Thus, the contents of the conditional-block can only be processed by the assembler one time.

The only code in our conditional-block is what is called a pseudo-instruction—the `EQU` command. The `EQU` command defines a symbol as equivalent to a constant value. In this case symbol `QW_SIZE` (which stands for: quadword size) is equated to the constant value 8. A quadword is 8-bytes in size.

NASM uses semicolons to indicate the beginning of a comment. A semicolon can begin at any position on a line. The semicolon and everything to the right of the semicolon is considered a comment and ignored by the NASM assembler.

The comment-block below defines our `memmove64` function. Enter the comment-block before the line with the `%endif` directive.

```

;-----
; C definition:
;
; void * memmove64 (void *dst, void const *src, size_t size)
;
; param:
;
; rdi = dst
; rsi = src
; rdx = size
;
; return:
;
; rax = dst
;
; WARNING: this routine does not handle the overlapping source-destination
;          senario.
;-----
;
%endif

```

The `memmove64` function is the workhorse of the list library. The comment-block provides: a C language declaration of function `memmove64`; a listing of the registers used to pass each parameter; the value `memmove64` will return (the address of the destination), and a warning about the constraints of the function.

Below we define the section of the object file our function will reside in—the `.text` section. We declare the symbol `memmove64` as `global` and of type `function`. The `global` directive works in conjunction with the `extern` directive that will be explained later. Should a shared library function be called during program execution, the dynamic linker will resolve the actual address of that function in a shared library and store that address in the Global Offset Table (GOT) for future reference. The Procedure Linkage Table (PLT) is used when a call is made from a shared library to a function outside that library. The PLT will access the address in the GOT once the dynamic linker has resolved the actual address. Next, the symbol `memmove64` is defined as a label, which is equivalent to the address of the function used during program execution.

Enter the following code-block below before the line with the `%endif` directive.

```
section .text
    global memmove64:function
memmove64:
    push    rdi                ; push destination address on stack
; quadword count = size / QW_SIZE
    mov     rax, rdx           ; copy dividend (size) to rax
    xor     rdx, rdx           ; zero out rdx
    mov     r11, QW_SIZE       ; copy divisor to r11
    div     r11                ; rdx:rax == (byte count):(quadword count)
    mov     rcx, rax           ; copy quadword count to rcx
    cld                        ; increment index registers rsi and rdi
    rep movsq                  ; repeat quadword move operation
    mov     rcx, rdx           ; copy byte count to rcx
    rep movsb                  ; repeat byte move operation
    pop     rax                ; pop destination address off stack
    ret
%endif
```

(Some functions will have a prologue and epilogue section. Such functions either have parameters that are passed on the stack (parameters seven and above) or have local variables. To access the parameters and/or local variables the `rbp` register will be used).

First, the contents of register `rdi` (the destination address for the move operation) is stored on the stack with the `push` instruction. The destination address is the return value of the function, and will be popped off the stack into register `rax` just before the function returns.

Second, the number of quadwords (the number of 8-byte blocks) is computed by dividing the value represented by `QW_SIZE` into the value passed in parameter `size`. In this case, the `div` instruction performs an unsigned division operation on a dividend in registers `rdx:rax` by a 64-bit divisor which can reside in memory or another register. (More on the `div` instruction can be found in the Intel 64

and IA32 Architecture Software Developer Manual). I use register `r11` to hold the divisor. The `div` instruction will place the quotient (quadword count) in register `rax` and the remainder (byte count) in register `rdx`.

Quotient = size / 8 = number times to call instruction `movsq` (move quadwords)

Remainder = size mod 8 = number of times to call instruction `movsb` (move bytes)

Third, the number of quadwords is copied from `rax` to `rcx` in order to be used by the repeat instruction `rep`. The `movsq` instruction moves a quadword (8-byte) string from the source address held in `rsi` to the destination address held in `rdi`. The instruction combo `rep movsq` moves 8-bytes per iteration, decrementing register `rcx` by one each iteration until the value in `rcx` reaches zero. Each iteration increments by eight the source address in register `rsi` and the destination address in register `rdi`.

Forth, the `direction-flag` is cleared (DF bit is 0) by the `cld` instruction. Clearing the `direction-flag` means setting the corresponding bit in the `eflags` register to zero. When the `direction-flag` is cleared the index registers `rsi` and `rdi` are incremented by string operations. When the `direction-flag` is set (DF bit is 1) the index registers are decremented by string operations.

Now, when the number of quadwords was calculated the remainder was put in register `rdx` by the `div` instruction. The remainder represents the number of bytes (less than 8) that remain to be moved. So, we copy the remainder to register `rcx` to prepare for the `rep movsb` instruction combo which will move the remaining bytes (if any) from the source address in `rsi` to the destination address in `rdi`, decrementing register `rcx` and incrementing registers `rsi` and `rdi` accordingly. (Notice that the only register to be modified between the two different move cycles was `rcx`. The source and destination registers both hold the correct addresses for the next iteration (if any) of 1-byte move operations).

The last operation before function `memmove64` returns is to set the return value in register `rax`. The return value is the destination address (the parameter passed in register `rdi`) that was pushed onto the stack prior to any operations that would have modified `rdi`. That address is popped off the stack into `rax` as the return value of function `memmove64`.

See Appendix B for the full source code of function `memmove64`.

LIST STRUCTURE IN ASSEMBLY

Now that we have our workhorse function we can build the structure and functions of our list library. Next, create a file `list.inc` in the `util/util/` via the command line or with your text editor.

Enter the code below into file `list.inc`.

```
; file:    list.inc
; brief:   list library structure definition
#ifndef list_INC
#define list_INC 1
;-----
```

```
;
%endif
```

We need a list structure definition in Assembly. Enter the following structure definition before the line with `%endif`.

```
struc list
    .o_size      resq    1      ; size (in bytes) of one object
    .total       resq    1      ; maximum # of objects this list can hold
    .count       resq    1      ; number of objects in list
    .index       resq    1      ; iterator index
    .buffer      resq    1      ; pointer to list buffer
endstruc
;
listSize      EQU      list_size
;
%endif
```

NASM provides the macros `struc` and `endstruc` to define a structure data type. The name of our structure is `list`. The structure is terminated by the `endstruc` macro. Each field in our list library uses the RESB family of psuedo-instructions in general and the RESQ (which means *reserve quadword*—equivalent to 64-bits or 8-bytes) in particular. As we work through coding the list library functions the purpose of each field will become clear. However, each field has a comment to the right that provides some insight into its purpose. All the fields are 64-bit unsigned integers with the exception of `.buffer` which is a 64-bit address (which is also a 64-bit unsigned integer).

To determine the size of a structure with NASM we append `_size` to our structure name `list` as in `list_size`. In this case, `struct list` is 40-bytes in size, so `list_size` equates to the integer value 40.

I define the symbol `listSize` as a constant equivalent to the value of `list_size` by using the `EQU` command.

LIST STRUCTURE IN C

We will need to define our list library in the C language as well. So, in the folder `util/util/` create the following files: `util.h` and `util.c`. Then open both files in your text editor. We will start with the C header file `util.h` by defining our list structure. Enter the following in file `util.h`.

```
#ifndef UTIL_H
#define UTIL_H

#include <stddef.h>
#include <stdint.h>
#include <stdlib.h>

typedef struct list list_t;
```

```

struct list {
    size_t    o_size;
    size_t    total;
    size_t    count;
    void *    index;
    void *    buffer;
};

```

This certainly seems clunky. Why two definitions for the same structure? The answer is simply that neither language understands the definition of the other. And each language needs a structure definition to work with. Furthermore, we will be using C to demo our Assembly functions, so it will be our C code that will instantiate the list structure and free the same.

We need `<stdint.h>` in order to use the C library definition for data type: `size_t`.

We need `<stdlib.h>` in order to use the C library memory allocation and de-allocation functions: `calloc` and `free`.

`size_t` on my computer is a 64-bit unsigned integer.

We use the same structure name and field names in C as in the Assembly code.

The keyword `typedef` is used to create an alias for a data type. In this case, `list_t` is another name for `struct list`.

In file `util.h` enter the following after the `list` structure definition.

```

};

#define list_alloc() (calloc(1, sizeof(list_t)))
#define list_free(P) (free(P))

```

The two defines above provide straight forward functions to allocate and free a `list` structure. For example, `list_alloc` will be expanded to `calloc(1, sizeof(struct list))` by the C pre-processor.

We will also need some C function declarations for our list library functions. To begin with we need a function that will initialize a list structure, and a function that will de-initialize the same. The initialization function needs some parameters, one of which is the address of a list structure and the other is the size (in bytes) of the objects we want to store in that structure.

Next, in file `util.h`, enter the following C function declarations below the allocation/de-allocation defines. Then enter the `#endif` directive to terminate the conditional-block.

```

#define list_free(P) (free(P))

```

```
int list_init (list_t *, size_t const);

void list_term (list_t *);

#endif
```

Next, in file `util.c` enter the following.

```
#include "util.h"
```

(I recommend leaving a blank line at the bottom of your `util.c` file. I might be wrong, but I seem to remember that being a good thing. Though it might not matter anymore).

LIST IMPLEMENTATION

At this point we need some Assembly code to go with these function declarations. In the folder `util/util/` create the file `list.asm`. Then, enter the following Assembly code into file `list.asm`.

```
; file:    list.asm
; brief:   list implementation
#ifdef LIST_ASM
#define LIST_ASM
;-----
extern bsearch
extern bzero
extern calloc
extern free
extern qsort
extern memmove64
;
#endif
```

As before we have a conditional-block of Assembly code. What is new is the **EXTERN** directive which declares a symbol as undefined in the module being processed by the assembler. However, the symbol is assumed to be defined and accessible elsewhere. In this case, `bsearch`, `bzero`, `calloc`, `free`, and `qsort` are C library (LIBC) functions, and `memmove64` is the function we wrote earlier. As for the LIBC functions, I will not cover them here. I recommend using the Linux man pages to learn more about them.

We need some constants to make our code more readable, so, enter the following into `list.asm` before the line with `%endif`.

```
;
LIST_COUNT      EQU      32
ALIGN_SIZE      EQU      16
ALIGN_WITH      EQU      (ALIGN_SIZE - 1)
ALIGN_MASK      EQU      ~(ALIGN_WITH)
```



```
;
%endif
```

The constant `LIST_COUNT` defines the number of slots available in our list library once it has been initialized by the function `list_init`. `ALIGN_SIZE`, `ALIGN_WITH` and `ALIGN_MASK` will be used to adjust our buffer size to an 16-byte boundary, as well as, adjust the stack when calling C callback and LIBC functions.

The `~` operator generates the bitwise negation (ones compliment) of its operand.

(`ALIGN_WITH`) is (`16 - 1`) or `15`. The 64-bit bitwise negation of `15` is:

```
11111111 11111111 11111111 11111111 11111111 11111111 11111111 11110000.
```

This value will be used to adjust our buffer size (and the stack) to a 16-byte boundary.

We will need a consistent way to align the stack to a 16-byte boundary and then call C functions. Enter the following code into `list.asm` before the line `%endif`.

```
;
%macro ALIGN_STACK_AND_CALL 2-4
    mov     %1, rsp           ; backup stack pointer (rsp)
    and     rsp, QWORD ALIGN_MASK ; align stack pointer (rsp) to
                                ; 16-byte boundary
    call    %2 %3 %4         ; call C function
    mov     rsp, %1          ; restore stack pointer (rsp)
%endmacro
;
%endif
```

The macro `ALIGN_STACK_AND_CALL` simplifies calls to some LIBC and C callback functions. The macro takes anywhere from two to four parameters as indicated by “2-4” after the macro name. We will pass two parameters for our C callback functions, since, the address of the callback is already known. But, for LIBC function calls we will pass four parameters. The macro stores the value of the Stack Point (register `rsp`) in a callee-saved register of our choice. Then, aligns the Stack Pointer down to a 16-byte boundary. Next the call is made to the function. Upon, return from the function the Stack Pointer is restored.

NOTE: We really could do without the `ALIGN_STACK_AND_CALL` macro for this tutorial. We won’t be dealing with single or double-precision values. However, if we were the macro would be necessary, for instance, the LIBC function `printf` requires the stack be aligned to a 16-byte boundary when passing single and double precision values. This is due to the use of Advanced Vector Extensions (AVX) instructions which we won’t be covering in this tutorial.

Our functions will need access to the Assembly definition of our list structure. Enter the following before the `%endif` directive.

```
;
#include "list.inc"
;
%endif
```

Now we declare a text section where the Assembly functions will reside. Enter the following before the `%endif` directive.

```
section .text
;
%endif
```

FUNCTION: LIST_INIT

The first Assembly function we will code is `list_init`. Enter the comment-block below into `list.asm` starting before the line with the `%endif`.

```
;
;-----
; C declaration:
;
;   int list_init (list_t *list, size_t const o_size);
;
; param:
;
;   rdi = list
;   rsi = o_size
;
; return:
;
;   eax = 0 (success) | -1 (failure)
;
; stack:
;
;   [rbp - 8]   = rdi (list)
;-----
;
%endif
```

The comments provide some insight into what we will be coding. For instance, the C declaration of function `list_init` is reproduced here. The parameters passed to `list_init` are passed via registers `rdi` and `rsi`. The `list_init` function returns a 32-bit integer result of either 0 for success or -1 for failure. And the value in register `rdi` (the address of the list structure) is stored on the stack. As you remember, parameters are passed to C functions on Linux in left-to-right order:

- parameters 1-6 are passed in registers `rdi`, `rsi`, `rdx`, `rcx`, `r8`, `r9` respectively.
- parameters above 6 are passed on the stack in reverse order (and can be popped off the stack in order).

Enter the code below into `list.asm` starting before the line with `%endif`.

```
        global list_init:function
list_init:
; prologue
        push     rbp
        mov      rbp, rsp
        sub      rsp, 8
;
; more code goes here
;
.epilogue:
        mov      rsp, rbp
        pop      rbp
        ret
%endif
```

The `global` directive is the partner of the `extern` directive. Whereas, `extern` assumes the symbol is declared elsewhere (in a separate module), the `global` directive provides that declaration. The symbol `list_init` is given type `function` by using an extension of the Executable and Linkable Format (ELF). Thus, the symbol `list_init` refers to a `function`. On the next line, the symbol is defined as a label which will become the offset to the beginning of function `list_init`. After that comes what I call the prologue of the function (some list library functions will have a prologue and an epilogue).

Register `rbp` is known as the Base (or Frame) Pointer Register and points to the base of the stack frame of the current function. However, as far as parameters and local variables are concerned the frame pointer `rbp` points to the center of the current frame. The `rbp` register is callee-saved meaning the called function must save the contents of the register before modifying it. Hence, the `push rbp` instruction in the prologue. Next, the instruction `mov rbp, rsp` copies the value of the stack pointer `rsp` into the base pointer `rbp`. Then register `rsp` is adjusted to make room for local variables (as indicated in the comments above). Hence, register `rbp` points to the end of the parameter portion of the frame and the beginning of the local variable portion of the frame.

The epilogue of the function reverses the affects of the prologue by restoring the stack pointer and base (or frame) pointer values of the calling function. The `ret` instruction updates the instruction pointer `rip`, thereby, continuing execution at the instruction following the call to `list_init` in the caller.

Enter the following code after the `sub` instruction.

```
        sub      rsp, 8
; store rdi (list) on stack
        mov      QWORD [rbp - 8], rdi
; list->o_size = o_size
        mov      QWORD [rdi + list.o_size], rsi
```

The first thing after the prologue is to store the address of the list structure on the stack. `QWORD` is a size specifier though the assembler will know that `rdi` is a 64-bit register and generate the appropriate code. However, the keyword `STRICT` can be used to force the size. Next the object size is copied into the corresponding member (`o_size`) of the list structure. To be clear, the object size is the size of the objects that will be stored in the list. The object size is used to calculate the list buffer size and is used by almost all the list library functions to work with objects in the list. The list buffer size is then adjusted to an 16-byte boundary. The adjustment is more an exercise in itself than something we need to worry about.

Enter the following code after the `mov` instruction.

```
        mov     QWORD [rdi + list.o_size], rsi
; buffer_size = o_size * LIST_COUNT
        mov     rax, rsi
        mov     rcx, QWORD LIST_COUNT
        mul     rcx
; buffer_size = (buffer_size + ALIGN_WITH) & ALIGN_MASK
        add     rax, QWORD ALIGN_WITH
        and     rax, QWORD ALIGN_MASK
```

The size of objects to be stored in the list is passed to `list_init` via register `rsi`. We will use that value to calculate the initial size of the list buffer. The first `mov` instruction copies the object size from register `rsi` to register `rax`. Then, `LIST_COUNT` is copied to register `rcx`. Next, the object size in register `rax` is multiplied by the value `LIST_COUNT` in register `rcx` and the result stored in the register combo `rdx:rax` by the `mul` instruction. (NOTE: I assume register `rdx` to be zero after the `mul` instruction). Then, the buffer size is adjusted to an 16-byte boundary by adding the value of `ALIGN_WITH` to the buffer size in register `rax` (the `add` instruction stores the result in the destination register which is the first of the two operands). Finally, the value in `rax` is aligned to an 16-byte boundary via the bitwise `and` instruction against the alignment mask `ALIGN_MASK`.

Enter the following code after the `and` instruction.

```
        and     rax, QWORD ALIGN_MASK
; if ((list->buffer = calloc(1, buffer_size)) == NULL) return -1
        mov     rdi, 1
        mov     rsi, rax
        call    calloc wrt ..plt
        mov     rdi, QWORD [rbp - 8]
        mov     QWORD [rdi + list.buffer], rax
        test    rax, rax
        jnz     .continue
        mov     eax, -1 ; return -1 (failure)
        jmp     .epilogue
.continue:
```

Now we will allocate the list buffer on the heap. To allocate the buffer we use the C library function `calloc`. The `calloc` function takes two arguments, the first is the number of objects and the second is the size of each object. With these values `calloc` calculates the size of the buffer to allocate. In our case, the size has already been determined. So, we pass `1` as the first argument in register `rdi`, and the calculated buffer size as the second argument in register `rsi`. Then a call is made to `calloc` with the result being returned in register `rax`. Next, the value returned from `calloc` is stored in the `buffer` member of the list structure. In order to access the list structure, we must restore its address from the stack to register `rdi`. We do this for two reasons: first, register `rdi` held the first parameter of function `calloc`; and second, register `rdi` is not a callee saved register. After storing the address of the buffer we check to see if the `calloc` function was successful. To do that we call the `test` instruction on register `rax` and check to determine if the `zero-flag` is *not* set with the `jnz` instruction. If the `zero-flag` was set by the `test` instruction then the `calloc` function returned a `NULL` pointer indicating failure, so, execution moves to the next instruction where register `eax` is set to `-1`, followed by a jump to the first instruction after label `.epilogue`, which leads to the termination of function `list_init`. Otherwise, function `calloc` returned a valid memory address, and we jump to the next instruction after the label `.continue`.

Enter the following code after label `.continue`.

```
.continue
; list->total = LIST_COUNT
    mov     rax, LIST_COUNT
    mov     QWORD [rdi + list.total], rax
; list->count = 0
    xor     rax, rax
    mov     QWORD [rdi + list.count], rax
; return 0
;     xor     eax, eax
```

At this point we have successfully allocated a buffer on the heap and can continue initializing the list structure. We set the `total` member of the list structure to the value of `LIST_COUNT` to indicate how many objects the list buffer can hold. To set the `total` member we use register `rdi` (which holds the address of the list structure) and add the offset to the member `total`. We use the structure we defined in file `list.inc` to provide the offset. The list structure amounts to a group of offsets we can use to access members in an actual list structure. Next we use the `xor` instruction to zero out the `rax` register. Then we use `rax` to set member `count` to zero. The last `xor` is commented out since `rax` has already been set to zero, so, we already have our return value.

See Appendix C for the full source code of function `list_init`.

FUNCTION: LIST_TERM

We need a function that will prepare a list structure to be deallocated. I call the function a list structure terminator. The list structure terminator is not that complicated, we just need to free the list buffer memory and zero out the list structure members. `list_term` will be that function.

Enter the comment-block below into `list.asm` before the `%endif`.

```
;
;-----
; C definition:
;
; void list_term (list_t *list)
;
; param:
;
; rdi = list
;
; stack:
;
; [rbp - 8] = rdi (list)
;-----
;
%endif
```

Unlike function `list_init` the `list_term` function has no return value. Hence, the `void` keyword before the function name in the C definition.

Enter the code below into `list.asm` before `%endif`.

```
global list_term:function
list_term:
; prologue
    push    rbp
    mov     rbp, rsp
    sub     rsp, 8
; store rdi (list) on stack
    mov     QWORD [rbp - 8], rdi
; free list buffer memory on heap
    mov     rdi, QWORD [rdi + list.buffer]
    call    free wrt ..plt
; zero out list structure
    mov     rdi, QWORD [rbp - 8]
    mov     rsi, QWORD listSize
    call    bzero wrt ..plt
; epilogue
    mov     rsp, rbp
    pop     rbp
    ret
%endif
```

`list_term` begins the same as the `list_init` function. The symbol `list_term` is declared as `global` and given the type `function`. The symbol is defined as a label, followed by the prologue of the function. The value in register `rdi` (the address of the list structure) is stored on the stack. Then the address of the list buffer is copied into register `rdi` as the first and only argument to the C library function `free`, which deallocates the corresponding memory on the heap. Next the address of the list structure is restored to register `rdi` as the first of two arguments to C library function `bzero`. The second argument, which is copied to register `rsi`, is the size of the list library (which is defined in file `list.inc`). `bzero` will zero out the entire list library and return. After that comes the epilogue of function `list_term`.

(NOTE: When making calls to functions in other shared libraries or the main program, NASM provides the keywords `wrt .plt` which stands for “With Reference To” “Procedure Linkage Table.” More information is available in the NASM documentation in sections “SEG and WRT” and “Calling Procedures Outside the Library”).

See Appendix D for the full source code of function `list_term`.

SHARED LIBRARY MAKEFILE

Our goal here is to create a shared library. We need to create a `makefile` that will do the work of creating the shared object file `libutil.so`. So, in folder `util/util/` create the file `makefile`.

Now enter the following into file `makefile`. (NOTE: use tabs instead of spaces to indent lines).

```
# makefile for libutil.so
libutil.so: memmove64.o list.o util.o
    gcc -z noexecstack -shared memmove64.o list.o \
        util.o -o libutil.so
util.o: util.c
    gcc -fPIC -c util.c -o util.o
list.o: list.asm
    nasm -f elf64 list.asm -o list.o
memmove64.o: memmove64.asm
    nasm -f elf64 memmove64.asm -o memmove64.o
clean:
    rm -f libutil.so util.o list.o memmove64.o
```

Now that we have a `makefile` we can build our shared object file `libutil.so`. At the command line change directory to folder `util/util/`. Now run the following at the command line.

```
make clean; make
```

Assuming everything went as expected you should see the following output.

```
rm -f libutil.so util.o list.o memmove64.o
nasm -f elf64 memmove64.asm -o memmove64.o
nasm -f elf64 list.asm -o list.o
gcc -fPIC -c util.c -o util.o
gcc -z noexecstack -shared memmove64.o list.o util.o -o libutil.so
```

The `clean` argument tells `make` to execute the commands in the `clean:` section of the `makefile`. The next call to `make` with no argument tells `make` to execute the commands that comprise the recipe that produces the prerequisites (in this case files `memmove64.o`, `list.o`, `util.o`) necessary to create the target `libutil.so`; and then, make the shared library `libutil.so`.

At the command line change directory to folder `util/util/` and enter the following command:

```
ls -C1 *.o *.so
```

The `ls` command lists the object files created by the `make` command in folder `util/util/`. If all went well you should see the following files.

```
libutil.so
list.o
memmove64.o
util.o
```

If you do not see all these files or encountered an error message during compilation you will need to check your source files and/or `makefile` for accuracy.

TEST PROGRAM: LIST_INIT

At the command line change directory to folder `util/list_init/` and run the following command combo:

```
make clean; make
```

Now, enter the following command:

```
./list_init
```

You should see the following output from test program `list_init` indicating that all went well.

TEST LIST_INIT AND LIST_TERM

LIST_INIT SUCCESSFUL

```
list:      total:      24
          count:      0
```



```
o_size:    48
buffer:    0xf332e0
```

LIST_TERM SUCCESSFUL

```
list:      total:    0
           count:    0
           o_size:    0
           buffer:    (nil)
```

The members of the `list` structure are shown after the `list` structure is initialized successfully and after it is de-initialized successfully. Take time to look at the source code for test program `list_init` in files `main.h` and `main.c` in folder `util/list_init/`. The `list_alloc` and `list_free` defines (found in file `util.h` in the folder `util/util/`) are also used.

FUNCTION: LIST_ADD

Now we need a function that will add stuff to our list. We will name the function `list_add`. The `list_add` function will place the new object after the last object in the list. This function will increment the `count` member of the list structure after adding an object. If the list is full the `list_add` function will return `-1` indicating failure. Otherwise, zero will be returned.

In folder `util/util/`, in file `util.h`, enter the following function declaration before the `#endif` directive.

```
int list_add (list_t *, void const *);
```

`%endif`

Now open file `list.asm` in folder `util/util/` in your text editor. Next, enter the comment-block below into `list.asm` before the line with `%endif`.

```
;
;-----
; C definition:
;
; void * list_add (list_t *list, void const *object);
;
; param:
;
; rdi = list
; rsi = object
;
; return:
;
; rax = address of object in list | NULL
;
; stack:
```

```

;
; [rbp - 8] = rdi (list)
; [rbp - 16] = (void *addr) = address of object in list buffer
;-----
;
%endif

```

Again we have the C declaration for function `list_add`; followed by the parameter `list` indicating which registers hold what data; then the possible return values are indicated; and finally the disposition of the stack.

Enter the following code-block into file `list.asm` before the line with `%endif`.

```

    global list_add:function
list_add:
; prologue
    push    rbp
    mov     rbp, rsp
    sub     rsp, 16
; store rdi (list) on stack
    mov     QWORD [rbp - 8], rdi
%endif

```

We have the function declaration as `global` with type `function` followed by the label definition and the prologue. The Stack Pointer (`rsp`) is adjusted by 16-bytes to provide storage for the address of the list structure passed in parameter `list`, and the address where, in the list, the new object will be stored. Then, the address of the list structure is stored on the stack.

Enter the following code-block into file `list.asm` before the line with `%endif`.

```

; if (list->count >= list->total) return NULL
    xor     rax, rax
    mov     rcx, QWORD [rdi + list.count]
    cmp     rcx, QWORD [rdi + list.total]
    jae     .epilogue
%endif

```

If the list buffer is full, then `NULL` is returned and the function terminates. Why zero out register `rax` with the `xor` instruction before the comparison is performed? By setting `rax` ahead of the `cmp` instruction only one jump instruction is necessary and the jump instruction mirrors the comparison in the comment above it. If the value in member `count` is above or equal-to the value in member `total` a jump will be made to the instruction following label `.epilogue`. The `jae` (jump, if above or equal-to) instruction is for unsigned comparisons only. For a signed comparison we would have used the `jge` (jump, if greater-than or equal-to) instruction.

Enter the following code-block into file `list.asm` before the line with `%endif`.

```

; void *slot = &list->buffer[(list->count * list->o_size)]
    mov     rax, QWORD [rdi + list.count]
    mul     QWORD [rdi + list.o_size]
    add     rax, QWORD [rdi + list.buffer]
    mov     QWORD [rbp - 16], rax
%endif

```

If there is room for another object in the list, then we need the address for that slot. That address is calculated by the first three lines of code above. The value in member `count` is copied to register `rax`. Then the count is multiplied by the object size in member `o_size` leaving the result in `rdx:rax`. (I ignore the value in register `rdx`, since my computer has nowhere near that much memory. Therefore, I expect the value in `rdx` to always be zero). The value in `rax` is an offset from the beginning of the list buffer. So, we will add the address in member `buffer` to the offset in `rax`. Now we have the address of the free slot in the list buffer. The last operation stores the address of the free slot on the stack.

Enter the following code-block into file `list.asm` before the line with `%endif`.

```

; (void) memmove64(slot, object, list->o_size)
    mov     rdx, QWORD [rdi + list.o_size]
    mov     rdi, rax
    call    memmove64 wrt ..plt
%endif

```

Now we move the new object into the list buffer. First, the size of an object in bytes is copied from member `o_size` into register `rdx`. Next, the address of the free slot is copied from `rax` to `rdi`. You may have noticed that we have not modified register `rsi` to this point. Therefore, `rsi` still contains the address of the `object` passed by the caller. So, we call function `memmove64` to copy the contents of the new object to the free slot in the list buffer.

Enter the following code-block into file `list.asm` before the line with `%endif`.

```

; list->count += 1
    mov     rdi, QWORD [rbp - 8]
    mov     rax, QWORD [rdi + list.count]
    inc     rax
    mov     QWORD [rdi + list.count], rax
; return addr
    mov     rax, QWORD [rbp - 16]
%endif

```

There are two operations to perform before the function epilogue. First, the value in member `count` must be incremented. To do that we restore the address of the list structure to register `rdi`; copy the value in member `count` to register `rax`; increment the value in `rax` by one with the `inc` instruction; and copy the incremented value into member `count`. (Incrementing the count could have been

achieved by incrementing the memory location directly). Second, we need to set the function return value in register `rax`. So, the address of the free slot, where the `object` was copied to, is copied to register `rax`.

Enter the following code-block into file `list.asm` before the line with `%endif`.

```
.epilogue:
    mov     rsp, rbp
    pop     rbp
    ret
%endif
```

The epilogue restores the Stack Pointer Register (`rsp`) and the Base (or Frame) Pointer Register (`rbp`), erasing any modifications made during execution of the `list_add` function. The `ret` instruction returns execution to the caller.

See Appendix E for the full source code of function `list_add`.

TEST PROGRAM: LIST_ADD

Since, we modified file `list.asm`, we will need to rebuild the shared library file `libutil.so`. At the command line change directory to folder `util/util/`, and run the following command combo:

```
make clean; make
```

Now, change directory to folder `util/list_add/`, and run the same command combo:

```
make clean; make
```

Now, run the `list_add` test program with the following command:

```
./list_add
```

You should see the following output from test program `list_add` indicating that all went well. (NOTE: the buffer address will differ on your system).

```
TEST LIST_ADD
```

```
LIST_INIT SUCCESSFUL
```

```
list:      total:      36
           count:      0
           o_size:     36
           buffer:     0x19402e0
```

```
list_add:  address: 0x19402e0    offset:  0
```

list_add:	address: 0x1940304	offset: 36
list_add:	address: 0x1940328	offset: 72
list_add:	address: 0x194034c	offset: 108
list_add:	address: 0x1940370	offset: 144
list_add:	address: 0x1940394	offset: 180
list_add:	address: 0x19403b8	offset: 216
list_add:	address: 0x19403dc	offset: 252
list_add:	address: 0x1940400	offset: 288
list_add:	address: 0x1940424	offset: 324
list_add:	address: 0x1940448	offset: 360
list_add:	address: 0x194046c	offset: 396
list_add:	address: 0x1940490	offset: 432
list_add:	address: 0x19404b4	offset: 468
list_add:	address: 0x19404d8	offset: 504
list_add:	address: 0x19404fc	offset: 540
list_add:	address: 0x1940520	offset: 576
list_add:	address: 0x1940544	offset: 612
list_add:	address: 0x1940568	offset: 648
list_add:	address: 0x194058c	offset: 684
list_add:	address: 0x19405b0	offset: 720
list_add:	address: 0x19405d4	offset: 756
list_add:	address: 0x19405f8	offset: 792
list_add:	address: 0x194061c	offset: 828

LIST_ADD SUCCESSFUL

list:	total:	36
	count:	24
	o_size:	36
	buffer:	0x19402e0

LIST_ADD FAILED AS EXPECTED!

list:	total:	24
	count:	24
	o_size:	36
	buffer:	0x19402e0

LIST_TERM SUCCESSFUL

list:	total:	0
	count:	0
	o_size:	0
	buffer:	(nil)

After each function or loop completes the members of the list structure are printed out. After `list_init` completes successfully the list structure print out indicates that: the total number of objects the list can hold is 24; the number of objects in the list is 0; the size of each object is 36-bytes; and the address of the buffer where the objects will be stored—`0x19402e0`. (NOTE: the address you see on your system will be different). Next, the contents of 24 vehicle structures are added to the list via a for-loop. The function name is printed along with the address of the free slot and the offset of the free slot from the beginning of the list buffer. After the for-loop completes the members of the list structure are displayed. This time the count indicates that 24 objects are in the list. The list is now full. To test the integrity of the function, one more call is made to `list_add`, and as expected the `list_add` function returns `NULL` indicating failure. Another display of the list structure confirms that no additional object was added. Next the `list_term` and `list_free` functions are called and the test program terminates. Take time to look at the source code for test program `list_add` in files `main.h` and `main.c` in folder `util/list_add/`.

FUNCTION: LIST_COUNT

This function is very simple and demonstrates how easy it is to access a member of a structure via an Assembly function. (Gives the list structure an object-oriented touch.)

In folder `util/util/`, in file `util.h`, enter the following function declaration before the `#endif` directive.

```
size_t list_count (list_t *);
```

```
%endif
```

Now open file `list.asm` in folder `util/util/` in your text editor. Next, enter the comment-block below into `list.asm` before the line with `%endif`.

```
;
;-----
; C definition:
;
;   size_t list_count (list_t *list);
;
; param:
;
;   rdi = list
;
; return:
;
;   rax = value in member list->count
;-----
;
%endif
```

Again we have the C declaration of function `list_count`; followed by the parameter list indicating one parameter is being passed via register `rdi`; and, then the return value is indicated.

Enter the following code-block into file `list.asm` before the line with `%endif`.

```
    global list_count:function
list_count:
    mov     rax, QWORD [rdi + list.count]
    ret
```

We have the function declaration as `global` with type `function` followed by the label definition. One parameter is passed to the function which is the address of a list structure. The value in the `count` member of the list structure is copied to register `rax`. In this case, the `rax` register is where the caller will look for the return value of this function. The `ret` instruction returns execution to the caller.

See Appendix F for the full source code of function `list_count`.

FUNCTION: LIST_BEGIN

Eventually we will want to iterate through the stuff we add to the list structure. To do that we need a point of reference—somewhere to begin the iteration. That is the purpose of function `list_begin`. Assuming there is stuff in the list buffer to iterate through, the function will provide an address to the first object in the list buffer. However, we should expect a NULL pointer, if the list buffer is empty.

In folder `util/util/`, in file `util.h`, enter the following function declaration before the `#endif` directive.

```
void * list_begin (list_t *list);
```

```
%endif
```

Now in file `list.asm` in folder `util/util/`, enter the comment-block below before the line with `%endif`.

```
;
;-----
; C definition:
;
; void * list_begin (list_t *list);
;
; param:
;
; rdi = list
;
; return:
```

```

;
;   rax = list->buffer | NULL
;-----
;
%endif

```

The comment block gives us the usual information. Enter the following code-block into file `list.asm` before the line with `%endif`.

```

    global list_begin:function
list_begin:
%endif

```

The symbol `list_begin` is declared `global` with type `function`, and then defined as a label. Next, enter the following code-block into file `list.asm` before the line with `%endif`.

```

; if (list->count == 0) return NULL
    mov     rax, QWORD [rdi + list.count]
    test    rax, rax
    jz      .epilogue
%endif

```

The comment line tells us what to expect from the Assembly code that follows. If the list structure is empty (indicated by the value of the `count` member being zero, then the function returns a `NULL` pointer. However, the code is a bit tricky. If the value of the `count` member is zero then zero is copied into register `rax`. The `test` instruction will set the `zero-flag` should the value in `rax` be zero. That is the requirement of the jump instruction `jz` which performs a jump to the `ret` instruction following the label `.epilogue`. The trick is killing to birds with one stone by using the zero in member `count` as the `NULL` pointer in register `rax`. As indicated in the comment block above, the function return value will be in register `rax`. The `rax` register is where the caller of this function expects to find the return value. Next, enter the following code-block into file `list.asm` before the line with `%endif`.

```

; list->index = 0L
    xor     rax, rax
    mov     QWORD [rdi + list.index], rax
%endif

```

If the value in the `count` member is greater than zero, then the `jz` instruction falls through and execution continues with setting the value of member `index` to zero. This is achieved by using the `xor` instruction to zero out the `rax` register. Then, the value in `rax` is copied into member `index`. Next, enter the following code-block into file `list.asm` before the line with `%endif`.

```

; return list->buffer
    mov     rax, QWORD [rdi + list.buffer]

```



```
.epilogue:
    ret
%endif
```

The last operation before the `ret` instruction is to copy the address from the `buffer` member to register `rax` as the return value of function `list_begin`.

See Appendix G for the full source code of function `list_begin`.

FUNCTION: LIST_NEXT

The `list_next` function is an incremental iterator, and works in conjunction with function `list_begin`.

In folder `util/util/`, in file `util.h`, enter the following function declaration before the `#endif` directive.

```
void * list_next (list_t *);

%endif
```

Next, enter the comment-block below into file `list.asm` before the line with `%endif`.

```
;
;-----
; C definition:
;
; void * list_next (list_t *list);
;
; param:
;
; rdi = list
;
; return:
;
; rax = &list->buffer[list->index] | NULL
;-----
;
%endif
```

Enter the following code-block into file `list.asm` before the line with `%endif`.

```
global list_next:function
list_next:
; if (list->index >= (list->count - 1)) return NULL
    xor     rax, rax
```

```

        mov     rcx, QWORD [rdi + list.count]
        dec     rcx
        mov     rdx, QWORD [rdi + list.index]
        cmp     rdx, rcx
        jae     .return
%endif

```

We have the usual function declaration and definition. The if statement on the comment-line has a comparison to determine whether or not the iterator has reached the end of the objects in the list. If so, then a NULL pointer will be returned. The offset (in bytes) of the beginning of the first object in the list buffer is zero. That is, the memory address of the list buffer is also the address of the first object in the buffer. Since, the offset of the first object in the buffer is 0 as illustrated below.

A	B	C	D	E	F
0	1	2	3	4	5

The letters represent objects in the list buffer and the numbers represent the corresponding index of each object.

The index of an object multiplied by the size of the objects in the list, results in the offset of that object from the beginning of the list. And, by adding the offset of an object to the address of the beginning of the list buffer you get the address of the object itself.

Since the index of the first object in the list is zero the corresponding offset will be zero, as the integer value zero multiplied by another integer value is always zero. This leaves the address of the list buffer as the address of the first object in the list. (So simple, and yet, so many words).

Enter the following code-block into file `list.asm` before the line with `%endif`.

```

; list->index += 1
    inc     rdx
    mov     QWORD [rdi + list.index], rdx
%endif

```

The purpose of function `list_next` is to return the address of the next object in the list buffer. Therefore, before that address can be calculated the value of the `index` member needs to be incremented. The value of member `index` is already in register `rdx`. So, the value in `rdx` is incremented; then copied into structure member `index`. (NOTE: we could have incremented the value in the `index` member directly with the following instruction:

```
inc     QWORD [rdi + list.index]
```

In this case, the size prefix `QWORD` is necessary, and indicates to the assembler that the memory operand is a 64-bit value).

Enter the following code-block into file `list.asm` before the line with `%endif`.

```

; return (list->buffer + (list->index * list->o_size))
    mov     rax, rdx
    mul     QWORD [rdi + list.o_size]
    add     rax, QWORD [rdi + list.buffer]
.return:
    ret
%endif

```

The code-block above, calculates the address of the next object in the list, and places that address in register `rax` as the return value. The value of the incremented `index` member in `rdx` is copied to register `rax`. An offset from the beginning of the list buffer is calculated by multiplying the value in `rax` by the value in structure member `o_size`. Finally, the address (of the next object in the list buffer) is calculated by adding the address of the buffer held in structure member `buffer` to the calculated offset held in register `rax`. Once the calculation is complete the address of the next object in the list resides in register `rax`, and the function returns via the `ret` instruction.

See Appendix H for the full source code of function `list_next`.

TEST PROGRAM: LIST_BEGIN

Since, we modified file `list.asm`, we will need to rebuild the shared library file `libutil.so`. At the command line change directory to folder `util/util/`, and run the following command combo:

```
make clean; make
```

Now, change directory to folder `util/list_begin/`, and run the same command combo:

```
make clean; make
```

Now, run the `list_begin` test program with the following command:

```
./list_begin
```

You should see the following output from test program `list_begin` indicating that all went well. (NOTE: the buffer address will differ on your system).

TEST LIST_BEGIN AND LIST_NEXT

LIST_INIT SUCCESSFUL

```

list:      total:      24
           count:      0
           o_size:     36
           buffer:     0x11f72e0

```

LIST_ADD SUCCESSFUL

```
list:      total:      24
           count:      24
           o_size:     36
           buffer:     0x11f72e0
```

01:	Ford	Aspire	1994
02:	Chevrolet	Silverado 1500	2003
03:	Buick	Skylark	1997
04:	BMW	Z4	2012
05:	BMW	Z4	2008
06:	Oldsmobile	Bravada	2002
07:	Pontiac	Grand Prix	1968
08:	Subaru	Legacy	2007
09:	GMC	Yukon	2006
10:	Mitsubishi	Truck	1991
11:	Hyundai	Tiburon	2001
12:	Dodge	Ram 3500	2002
13:	Alfa Romeo	164	1993
14:	GMC	Yukon XL 2500	2005
15:	Lexus	LS	1994
16:	Audi	Q7	2007
17:	Ford	Explorer	2007
18:	Pontiac	Grand Prix	1987
19:	Mercury	Capri	1993
20:	Hyundai	Equus	2012
21:	GMC	Savana 2500	2005
22:	Lexus	RX	2010
23:	Lexus	SC	2007
24:	Volkswagen	Cabriolet	1999

LIST_BEGIN AND LIST_NEXT SUCCESSFUL

LIST_TERM SUCCESSFUL

```
list:      total:      0
           count:      0
           o_size:     0
           buffer:     (nil)
```

After each function (or iteration of functions) completes the members of the list structure are printed out. After `list_init` completes successfully the list structure print out indicates that: the total number of objects the list can hold is 24; the number of objects in the list is 0; the size of each object is 36-bytes; and the address of the buffer where the objects will be stored—0x11f72e0. (NOTE: the

address you see on your system will be different). Next, the contents of 24 vehicle structures are added to the list via a for-loop. After the for-loop completes the members of the `list` structure are displayed. This time the count indicates that 24 objects are in the list. The list is now full. The contents of 24 vehicle structures in the list buffer are displayed via a for-loop. Next, the `list_term` and `list_free` functions are called and the test program terminates. Take time to look at the source code for test program `list_begin` in files `main.h` and `main.c` in folder `util/list_begin/`. The iteration functions `list_begin` and `list_next` are used in a for-loop to iterate through the list of vehicle structures.

FUNCTION: LIST_SORT

We need a way to sort our list of vehicles by make, model, year or some combination of the three. In folder `util/util/`, in file `util.h`, enter the following function declaration before the `#endif` directive.

```
void list_sort (list_t *,
               int (*sort_cb) (void const *, void const *));
```

```
%endif
```

Next, enter the comment-block below into `list.asm` before the line with `%endif`.

```
;
;-----
; C definition:
;
; void list_sort (list_t *list, int (*sort_cb) (void const *, void const *));
;
; param:
;
; rdi = list
; rsi = sort_cb
;-----
;
%endif
```

The sort function takes two parameters. The first is the address of a `list` structure, and the second is the address of a callback function (`sort_cb`). The callback function takes two parameters, each of which is an object in the list buffer. The callback function is a compare function provided by the caller that must return a 32-bit integer value that is less-than, equal-to, or greater-than zero. The callback functions will be C functions and part of the test program.

Enter the following code-block into file `list.asm` before the line with `%endif`.

```
global list_sort:function
list_sort:
```

```

    push    r15
    mov     rcx, rsi
    mov     rdx, QWORD [rdi + list.o_size]
    mov     rsi, QWORD [rdi + list.count]
    mov     rdi, QWORD [rdi + list.buffer]
    ALIGN_STACK_AND_CALL r15, qsort, wrt, ..plt
    pop     r15
    ret

```

```
%endif
```

Yup, we be calling a C library function to do our sorting for us. The `qsort` function takes four parameters. We will copy the parameters in reverse order, though, all that matters is that each parameter is copied into the right register prior to the call to `qsort`. (If you need a refresher on what registers are involved in parameter passing, see the introduction). Register `rsi` holds the callback parameter (`sort_cb`) but will need to hold the `nmemb` parameter for the `qsort` function. For that reason I chose the order above. The `list_sort` function amounts to a wrapper function, but does simplify the call to `qsort` for the caller. We use the macro `ALIGN_STACK_AND_CALL` to make to the call to `qsort`. The macro sets the Stack Pointer to a 16-byte boundary. I use the macro for two reasons. First, it provides a consistent method of getting the stack pointer set correctly before the call. Second, the code is cleaner and the macro name informs us as to what we're doing. Register `r15` is a callee saved register (see introduction), and so we push the value in `r15` on the stack first thing. In the macro, the value in register `rsp` (the Stack Pointer) is stored in register `r15` before the call to `qsort`. After `qsort` returns the value in `r15` is copied back to register `rsp`, thus restoring the Stack Pointer. The macro will accept anywhere from 2 to 4 parameters, hence, the `2-4` in the macro definition. In this case, we use all four parameters, since, `qsort` is a LIBC function located somewhere outside this module.

See Appendix I for the full source code of function `list_sort`.

TEST PROGRAM: LIST_SORT

Since, we modified file `list.asm`, we will need to rebuild the shared library file `libutil.so`. So, at the command line change directory to folder `util/util/`, and run the following command combo:

```
make clean; make
```

Now, change directory to folder `util/list_sort/`, and run the same command combo:

```
make clean; make
```

Now, run the `list_sort` test program with the following command:

```
./list_sort
```

You should see the following output from test program `list_sort` indicating that all went well.
(NOTE: the buffer address will differ on your system).

TEST LIST_SORT

LIST_INIT SUCCESSFUL

```
list: total:    24
      count:     0
      o_size:   36
      buffer:   0x14712e0
```

LIST_ADD SUCCESSFUL

```
list: total:    24
      count:    24
      o_size:   36
      buffer:   0x14712e0
```

01:	GMC	Yukon	2006
02:	Ford	Aspire	1994
03:	Chevrolet	Silverado 1500	2003
04:	Buick	Skylark	1997
05:	BMW	Z4	2012
06:	BMW	Z4	2008
07:	Oldsmobile	Bravada	2002
08:	Pontiac	Grand Prix	1968
09:	Subaru	Legacy	2007
10:	Mitsubishi	Truck	1991
11:	Hyundai	Tiburon	2001
12:	Dodge	Ram 3500	2002
13:	Alfa Romeo	164	1993
14:	GMC	Yukon XL 2500	2005
15:	Lexus	LS	1994
16:	Audi	Q7	2007
17:	Ford	Explorer	2007
18:	Pontiac	Grand Prix	1987
19:	Mercury	Capri	1993
20:	Hyundai	Equus	2012
21:	Lexus	RX	2010
22:	Lexus	SC	2007
23:	Volkswagen	Cabriolet	1999
24:	GMC	Savana 2500	2005

RANDOM LIST SUCCESSFUL

01:	Alfa Romeo	164	1993
02:	Audi	Q7	2007

03:	BMW	Z4	2012
04:	BMW	Z4	2008
05:	Buick	Skylark	1997
06:	Chevrolet	Silverado 1500	2003
07:	Dodge	Ram 3500	2002
08:	Ford	Aspire	1994
09:	Ford	Explorer	2007
10:	GMC	Yukon	2006
11:	GMC	Yukon XL 2500	2005
12:	GMC	Savana 2500	2005
13:	Hyundai	Tiburon	2001
14:	Hyundai	Equus	2012
15:	Lexus	LS	1994
16:	Lexus	RX	2010
17:	Lexus	SC	2007
18:	Mercury	Capri	1993
19:	Mitsubishi	Truck	1991
20:	Oldsmobile	Bravada	2002
21:	Pontiac	Grand Prix	1968
22:	Pontiac	Grand Prix	1987
23:	Subaru	Legacy	2007
24:	Volkswagen	Cabriolet	1999

LIST_SORT BY MAKE SUCCESSFUL

01:	Alfa Romeo	164	1993
02:	Ford	Aspire	1994
03:	Oldsmobile	Bravada	2002
04:	Volkswagen	Cabriolet	1999
05:	Mercury	Capri	1993
06:	Hyundai	Equus	2012
07:	Ford	Explorer	2007
08:	Pontiac	Grand Prix	1968
09:	Pontiac	Grand Prix	1987
10:	Lexus	LS	1994
11:	Subaru	Legacy	2007
12:	Audi	Q7	2007
13:	Lexus	RX	2010
14:	Dodge	Ram 3500	2002
15:	Lexus	SC	2007
16:	GMC	Savana 2500	2005
17:	Chevrolet	Silverado 1500	2003
18:	Buick	Skylark	1997
19:	Hyundai	Tiburon	2001
20:	Mitsubishi	Truck	1991
21:	GMC	Yukon	2006
22:	GMC	Yukon XL 2500	2005
23:	BMW	Z4	2012

24: BMW

Z4

2008

LIST_SORT BY MODEL SUCCESSFUL

LIST_TERM SUCCESSFUL

```
list: total:      0
      count:      0
      o_size:     0
      buffer:     (nil)
```

The first listing of vehicles is random and in the order that the vehicle structures were added to the list. The second listing is sorted by vehicle make, which is the second column from the left. The third listing is sorted by vehicle model, which is the third column from the left. Take time to look at the source code for test program `list_sort` in files `main.h` and `main.c` in folder `util/list_sort/`.

FUNCTION: LIST_FIND

In order to find an object in our list structure we have two options: first, iterate through the list of objects until we find what we are looking for; second, sort the list and use a binary search to speed up the process. So, we need a binary search function. My solution is yet another wrapper function—`list_find`.

In folder `util/util/`, in file `util.h`, enter the following function declaration before the `%endif` directive.

```
void * list_find (list_t const *, void const *,
                 int (*find_cb) (void const *, void const *));
```

`%endif`

Now open file `list.asm` in folder `util/util/` in your text editor, and enter the comment-block below before the line with `%endif`.

```
;
;-----
; C declaration:
;
; void * list_find (list_t const *list, void const *key,
; int (*find_cb) (void const *, void const *));
;
; param:
;
; rdi = list
; rsi = key
; rdx = find_cb
```

```

;
; return:
;
;     eax = target (address of matching object) | NULL
;-----
;

```

The find function has three parameters: the address of a list structure; the address of a key to be searched for; and the address of a callback function. As before, the callback function is a compare function provided by the caller that must return a 32-bit integer value that is less-than, equal-to, or greater-than zero. The callback functions will be C functions found in the test program. Unlike function `list_sort`, the `list_find` function has a return value. The return value must be either an address of the found object or a `NULL` pointer.

Enter the following code-block into file `list.asm` before the line with `%endif`.

```

    global list_find:function
list_find:
    push    r12
    push    rsi
    mov     r8, rdx
    mov     rcx, QWORD [rdi + list.o_size]
    mov     rdx, QWORD [rdi + list.count]
    mov     rsi, QWORD [rdi + list.buffer]
    pop     rdi
    ALIGN_STACK_AND_CALL r12, bsearch, wrt, ..plt
    pop     r12
    ret

```

First, the contents of (callee-saved) register `r12` is pushed on the stack. The address of a key is passed in via register `rsi`. But, the `bsearch` function expects the address of the list buffer to be in register `rsi`. For that reason I chose to push the address in register `rsi` onto the stack. Next, the address of the callback function (parameter `find_cb`) is copied into register `r8`. (If you need a refresher on what registers are involved in parameter passing see the introduction). I chose to copy the parameters into the registers in reverse order (though all that matters is that the parameter be in the correct register for the call to `bsearch`). Next, the value in member `o_size` is copied to register `rcx`; then, the value in member `count` is copied to register `rdx`; next, the address in member `buffer` is copied to register `rsi`; and then, the address of a key to be searched for is popped off the stack into register `rdi`. Then, the macro `ALIGN_STACK_AND_CALL` does its thing and the call is made to `bsearch`. Upon completion, `bsearch` will return an address in register `rax` (or `NULL`). With the return value in register `rax`, nothing further need be done, but, to restore the contents of (callee-saved) register `r12` and return from the `list_find` function.

See Appendix J for the full source code of function `list_find`.

TEST PROGRAM: LIST_FIND

Since, we modified file `list.asm`, we will need to rebuild the shared library file `libutil.so`. So, at the command line change directory to folder `util/util/`, and run the following command combo:

```
make clean; make
```

Now, change directory to folder `util/list_find/`, and run the same command combo:

```
make clean; make
```

Now, run the `list_find` test program with the following command:

```
./list_find
```

You should see the following output from test program `list_find` indicating that all went well. (NOTE: the buffer address will differ on your system).

```
TEST LIST_SORT
```

```
LIST_INIT SUCCESSFUL
```

```
list: total:      24
      count:       0
      o_size:     36
      buffer:     0xedb2e0
```

```
LIST_ADD SUCCESSFUL
```

```
list: total:      24
      count:      24
      o_size:     36
      buffer:     0xedb2e0
```

01:	Alfa Romeo	164	1993
02:	Audi	Q7	2007
03:	BMW	Z4	2012
04:	BMW	Z4	2008
05:	Buick	Skylark	1997
06:	Chevrolet	Silverado 1500	2003
07:	Dodge	Ram 3500	2002
08:	Ford	Aspire	1994
09:	Ford	Explorer	2007
10:	GMC	Yukon	2006
11:	GMC	Yukon XL 2500	2005
12:	GMC	Savana 2500	2005
13:	Hyundai	Tiburon	2001
14:	Hyundai	Equus	2012

15:	Lexus	LS	1994
16:	Lexus	RX	2010
17:	Lexus	SC	2007
18:	Mercury	Capri	1993
19:	Mitsubishi	Truck	1991
20:	Oldsmobile	Bravada	2002
21:	Pontiac	Grand Prix	1968
22:	Pontiac	Grand Prix	1987
23:	Subaru	Legacy	2007
24:	Volkswagen	Cabriolet	1999

LIST_SORT BY VEHICLE MAKE SUCCESSFUL

LIST_FIND FOUND ALFA ROMEO

LIST_FIND FOUND LEXUS LS

LIST_FIND FOUND VOLKSWAGEN CABRIOLET

LIST_FIND FAILED TO FIND BUICK LUCERNE AS EXPECTED

LIST_TERM SUCCESSFUL

```
list: total:    0
      count:    0
      o_size:   0
      buffer:   (nil)
```

The listing of vehicles is sorted by vehicle make, which is the second column from the left. A sorted list, sorted on the search key, is necessary for a binary search. Then a search is made (with function `list_find`) for an Alfa Romeo. That search is successful. Next, a search is made for a Lexus LS, and that search is successful. Then, a search is made for a Volkswagen Cabriolet, which is also successful. Finally, a search is made for a Buick Lucerne, and that search fails as expected. Two of the searches were to test corner cases. For instance, the Alfa Romeo is the first vehicle in the list; and the Volkswagen Cabriolet is the last vehicle in the list. The Lexus LS is somewhere in the middle of the list. Finally, a search is made for a vehicle that is not in the list (the Buick Lucerne). Of course, no one expected the C library function `bsearch` to fail. Our tests have proven only that the `list_find` function and the supporting callback function `find_cb` were working correctly. Take time to look at the source code for test program `list_find` in files `main.h` and `main.c` in folder `util/list_find/`.

FUNCTION: LIST_DELETE

We have a way to add stuff to a list, so, how about a way to delete stuff. The `list_delete` function will remove an object from any slot in the list. The delete function will then shift objects left (if necessary) to fill the empty slot and keep the list contiguous. The function will decrement the `count` member of the list structure after removing an object from the list. And, return a value indicating the success or failure of the delete operation. For instance, if the list is empty the delete operation will fail

and the function will return `-1`. Or, if the address passed to the function is outside the address space of the list of objects in the list buffer, the function will fail and return `-1`.

In folder `util/util/`, in file `util.h`, enter the following function declaration before the `#endif` directive.

```
int list_delete (list_t *, void const *);

#endif
```

Now open file `list.asm` in folder `util/util/` in your text editor. Next, enter the comment-block below into `list.asm` before the line with `%endif`.

```
;
;-----
; C definition:
;
;   int list_delete (list_t *list, void const *key,
;       int (*find_cb) (void const *, void const *),
;       void (*delete_cb) (void const *));
;
; param:
;
;   rdi = list
;   rsi = key
;   rdx = find_cb
;   rcx = delete_cb
;
; return:
;
;   0 (success) | -1 (failure)
;
; stack:
;
;   QWORD [rbp - 8]   = rdi (list)
;   QWORD [rbp - 16]  = rcx (delete_cb)
;   QWORD [rbp - 24]  = (void *target)
;   QWORD [rbp - 32]  = (void *blk_tail)
;-----
;
#endif
```

There are four parameters and four local variables. The parameters include the address of a list structure; the address of a `key` (which could be any basic data type or structure); the address of a callback function for function `list_find`; and the address of a callback function for

`list_delete`. The `delete_cb` function provides a way for the user to de-initialize the `target` object prior to deletion. For instance, the `target` object may contain pointers to memory on the heap that need to be freed; or file descriptors that need closing; or any number of other reasons the user might want access to the `target` object prior to deletion.

Enter the following code-block into file `list.asm` before the line with `%endif`.

```
        global list_delete:function
list_delete:
; prologue
        push        rbp
        mov         rbp, rsp
        sub         rsp, 32
; store rdi (list) and rcx (delete_cb) on stack
        mov         QWORD [rbp - 8], rdi
        mov         QWORD [rbp - 16], rcx
%endif
```

The code-block includes the function declaration, label definition, and prologue of the `list_delete` function. The address of a list structure and the address of the `delete_cb` function are stored on the stack.

Now, enter the following code-block into file `list.asm` before the line with `%endif`.

```
; if ((target = list_find(list, key, find_cb)) == NULL) return -1
        call        list_find
        mov         QWORD [rbp - 24], rax
        test        rax, rax
        jnz         .target_found
        mov         eax, -1
        jmp         .epilogue
.target_found:
%endif
```

As the C code in the comment-line states, the `list_find` function is called with the `key` and `find_cb` function. This looks a bit strange, since, we did not provide the parameters to the `list_find` function. As it turns out, the first three parameters to `list_delete` are also the parameters to `list_find`. Those parameters are still in the appropriate registers, so, the call to `list_find` is made. If the `target` object is found, then the address of the `target` object will be in register `rax` upon return from function `list_find`. Otherwise, a `NULL` pointer is returned in `rax`. The `test` instruction will set the `zero-flag` should register `rax` hold a `NULL` value. If the `test` instruction does not set the `zero-flag`, then the `jnz` instruction (jump, if `zero-flag` not set) will perform a jump to label `.target_found`. However, should the `zero-flag` be set, then instruction `jnz` will fall through (not perform a jump) and execution will continue with copying `-1` to register `rax` to indicate that the `list_delete` function has failed. Next, the `jmp` instruction will

update the instruction pointer (register `rip`) with the address of the instruction following the `label.epilogue` where execution will continue.

Enter the following code-block into file `list.asm` before the line with `%endif`.

```
; delete_cb(target)
    mov     rcx, QWORD [rbp - 16]
    test    rcx, rcx
    jz      .no_delete_cb
    mov     rdi, rax
    call    rcx
.no_delete_cb:
%endif
```

If execution reaches the code-block above, then the `target` object has been found in the list buffer. In that case, the address of the `delete_cb` function is copied from the stack into register `rcx`. Followed by a test to determine if the address of the `delete_cb` function is `NULL` (the user can pass a `NULL` pointer, if the `delete_cb` function is not needed). As before, if the `delete_cb` address is a `NULL` pointer, then the `zero-flag` is set and the `jz` instruction (jump, if `zero-flag` set) will update the instruction pointer (`rip`) with the address of the instruction following label `.no_delete_cb`. Otherwise, the `jz` instruction will fall through and execution will continue with moving the address of the `target` object from register `rax` to register `rdi` (the first and only parameter to the `delete_cb` function). Until now, we have passed names to the `call` instruction. Now, we pass a register which holds the address of the function we want to call. So, let's take a moment and think this through. When we call a function by name what really takes place? Well, the assembler generates a call to the Procedure Linkage Table if the function is outside the module where the call took place. Otherwise, the address of the called function (a function within the module) is used. However, in this case, we provided the address of the function via a parameter to the `list_delete` function. And that address is in register `rcx`, which is a valid operand to the `call` instruction. (For more on the `call` instruction, see chapter 3 subsection 3 “INSTRUCTIONS (A-L)” of the “Intel 64 and IA-32 Architectures Software Developer’s Manual.”)

Now, enter the following code-block into file `list.asm` before the line with `%endif`.

```
; void *blk_tail = list->buffer + (list->count * list->o_size)
    mov     rdi, QWORD [rbp - 8]
    mov     rax, [rdi + list.count]
    mul     QWORD [rdi + list.o_size]
    add     rax, QWORD [rdi + list.buffer]
    mov     QWORD [rbp - 32], rax
%endif
```

The code-block above will calculate the address pointing to the end of the last object in the list buffer (`blk_tail`). First, we restore the address of the list structure (the first parameter to function `list_delete`) to register `rdi` from the stack. Then, we perform the equation in the comment line to calculate the address pointing to the end of the last object in the list buffer. We copy the value in

member `count` to register `rax`. Next, we multiply the value in register `rax` by the value in member `o_size`. The result of the multiplication is the size of the block of objects in the list buffer (think, adding the size of all the objects in the list together). Then, we add the address in member `buffer` to the calculated offset in register `rax`. The resulting address points to the end of the last object in the list buffer. We store on the address on the stack. (NOTE: The reason register `rax` is used for most math operations is due to the `mul` instruction. The implied first operand of the `mul` instruction is one of the following registers: `RAX`, `EAX`, `AX`, or `AL`. And the result of the `mul` instruction is stored in `AX` (`AH:AL`), `DX:AX`, `EDX:EAX`, or `RDX:RAX`).

Now, enter the following code-block into file `list.asm` before the line with `%endif`.

```
; if (target == (blk_tail - list->o_size)) goto .dec_count
    sub     rax, QWORD [rdi + list.o_size]
    cmp     QWORD [rbp - 24], rax
    je      .dec_count
%endif
```

Next, we need to determine whether or not the `target` object is the last object in the list buffer. The reason for this is that deleting the last object in the list amounts to decrementing the value in member `count` (no move operation is necessary). So, to start with, register `rax` holds the address that points to the end of the last object in the list. We will subtract the size of an object from that address to get the address of the last object in the list. With that done, we compare the address of the `target` object with the address of the last object in the list. Should the addresses be equivalent, execution jumps to the first instruction following label `.dec_count`.

Enter the following code-block into file `list.asm` before the line with `%endif`.

```
; void *blk_head = target + list->o_size
    mov     rcx, QWORD [rbp - 24]
    add     rcx, QWORD [rdi + list.o_size]
%endif
```

Now, having determined that the `target` address is not the address of the last object in the list buffer, execution proceeds with calculating the address of the object immediately following the `target` object in the `list` buffer. This address will be the `blk_head`, the address pointing to the beginning of the block of objects that we need to shift left one slot to fill the gap left by the deleted object (`target`). The process is simple enough. We add the value of member `o_size` to the address of the `target` object. The `target` address is copied to register `rcx`. Next, the value of member `o_size` (object size) is added to the value in register `rcx`. With the `blk_head` address calculated execution will proceed to the next block of code.

Enter the following code-block into file `list.asm` before the line with `%endif`.

```
; size_t blk_size = blk_tail - blk_head
    mov     rax, QWORD [rbp - 32]
    sub     rax, rcx
```


`%endif`

The `memmove64` function has three parameters: the destination address; the source address; and the size of the block of bytes to be copied. In this case, the destination address is the address of the `target` object; the source address is the `blk_head` address. So, all that remains to be calculated is the size (in bytes) of the block to be copied. The block size (`blk_size`) is calculated by subtracting the `blk_head` address from the `blk_tail` address. The `blk_tail` address was calculated earlier and stored on the stack. So, the `blk_tail` address is copied from the stack into register `rax` by the `mov` instruction. The `blk_head` address was just calculated and resides in register `rcx`. The subtraction is performed and the `sub` instruction puts the resulting value (`blk_size`) into the destination register `rax`.

Enter the following code-block into file `list.asm` before the line with `%endif`.

```
; memmove64(target, blk_head, blk_size)
    mov     rdi, QWORD [rbp - 24]
    mov     rsi, rcx
    mov     rdx, rax
    call    memmove64 wrt ..plt
%endif
```

With the three parameters calculated and placed in the correct registers, a call is made to the `memmove64` function to copy the block of objects one slot (object size) to the left.

Enter the following code-block into file `list.asm` before the line with `%endif`.

```
.dec_count:    ; list->count -= 1
    mov     rdi, QWORD [rbp - 8]
    mov     rax, QWORD [rdi + list.count]
    dec     rax
    mov     QWORD [rdi + list.count], rax
%endif
```

Another operation to be performed is adjusting the `count` member of the list structure to indicate the deletion of an object. Before this can be done the address of the list structure needs to be restored to register `rdi`. Next, the value in the `count` member of the list structure is copied to register `rax`. Then, the value in register `rax` is decremented by one. Next, the decremented `count` value is copied back to the `count` member in the list structure. (NOTE: The decrement operation could have been performed directly on the value in the list structure member `count`, as follows:

```
    dec     QWORD [rdi + list.count]
```

The size prefix `QWORD` is necessary, and indicates to the assembler that the memory operand is a 64-bit value.)

Enter the following code-block into file `list.asm` before the line with `%endif`.

```

; return 0
    xor     eax, eax
.epilogue:
    mov     rsp, rbp
    pop     rbp
    ret
%endif

```

With the count value updated we move on to setting the return value of the function (in this case zero), and performing the epilogue operations that essentially reverse what was done in the prologue. Then the `ret` instruction returns execution to the caller.

See Appendix I for the full source code of function `list_delete`.

TEST PROGRAM: LIST_DELETE

Since, we modified file `list.asm`, we will need to rebuild the shared library file `libutil.so`. At the command line change directory to folder `util/util/`, and run the following command combo:

```
make clean; make
```

Now, change directory to folder `util/list_delete/`, and run the same command combo:

```
make clean; make
```

Now, run the `list_delete` test program with the following command:

```
./list_delete
```

You should see the following output from test program `list_delete` indicating that all went well. (NOTE: the buffer address will differ on your system).

```
TEST LIST_DELETE
```

```
LIST_INIT SUCCESSFUL
```

```
list:      total:      24
           count:      0
           o_size:     36
           buffer:     0x1cf42e0
```

```
LIST_ADD SUCCESSFUL
```

```
list:      total:      24
           count:      24
```

o_size: 36
buffer: 0x1cf42e0

01:	GMC	Yukon	2006
02:	Ford	Aspire	1994
03:	Chevrolet	Silverado 1500	2003
04:	Buick	Skylark	1997
05:	BMW	Z4	2012
06:	BMW	Z4	2008
07:	Oldsmobile	Bravada	2002
08:	Pontiac	Grand Prix	1968
09:	Subaru	Legacy	2007
10:	Mitsubishi	Truck	1991
11:	Hyundai	Tiburon	2001
12:	Dodge	Ram 3500	2002
13:	Alfa Romeo	164	1993
14:	GMC	Yukon XL 2500	2005
15:	Lexus	LS	1994
16:	Audi	Q7	2007
17:	Ford	Explorer	2007
18:	Pontiac	Grand Prix	1987
19:	Mercury	Capri	1993
20:	Hyundai	Equus	2012
21:	Lexus	RX	2010
22:	Lexus	SC	2007
23:	Volkswagen	Cabriolet	1999
24:	GMC	Savana 2500	2005

LIST_BEGIN AND LIST_NEXT SUCCESSFUL

LIST_SORT SUCCESSFUL

01:	Alfa Romeo	164	1993
02:	Audi	Q7	2007
03:	BMW	Z4	2012
04:	BMW	Z4	2008
05:	Buick	Skylark	1997
06:	Chevrolet	Silverado 1500	2003
07:	Dodge	Ram 3500	2002
08:	Ford	Aspire	1994
09:	Ford	Explorer	2007
10:	GMC	Savana 2500	2005
11:	GMC	Yukon	2006
12:	GMC	Yukon XL 2500	2005

13:	Hyundai	Equus	2012
14:	Hyundai	Tiburon	2001
15:	Lexus	LS	1994
16:	Lexus	RX	2010
17:	Lexus	SC	2007
18:	Mercury	Capri	1993
19:	Mitsubishi	Truck	1991
20:	Oldsmobile	Bravada	2002
21:	Pontiac	Grand Prix	1968
22:	Pontiac	Grand Prix	1987
23:	Subaru	Legacy	2007
24:	Volkswagen	Cabriolet	1999

LIST_BEGIN AND LIST_NEXT SUCCESSFUL

DELETING GMC VEHICLES FROM LIST

delete_cb:	GMC	Savana 2500	2005
delete_cb:	GMC	Yukon XL 2500	2005
delete_cb:	GMC	Yukon	2006

LIST_DELETE SUCCESSFUL

list:	total:	24
	count:	21
	o_size:	36
	buffer:	0x1cf42e0

01:	Alfa Romeo	164	1993
02:	Audi	Q7	2007
03:	BMW	Z4	2012
04:	BMW	Z4	2008
05:	Buick	Skylark	1997
06:	Chevrolet	Silverado 1500	2003
07:	Dodge	Ram 3500	2002
08:	Ford	Aspire	1994
09:	Ford	Explorer	2007
10:	Hyundai	Equus	2012
11:	Hyundai	Tiburon	2001
12:	Lexus	LS	1994
13:	Lexus	RX	2010
14:	Lexus	SC	2007
15:	Mercury	Capri	1993
16:	Mitsubishi	Truck	1991

17:	Oldsmobile	Bravada	2002
18:	Pontiac	Grand Prix	1968
19:	Pontiac	Grand Prix	1987
20:	Subaru	Legacy	2007
21:	Volkswagen	Cabriolet	1999

LIST_BEGIN AND LIST_NEXT SUCCESSFUL

LIST_DELETE FAILED AS EXPECTED

list:	total:	24
	count:	21
	o_size:	36
	buffer:	0x1cf42e0

LIST_DELETE FAILED AS EXPECTED

list:	total:	24
	count:	21
	o_size:	36
	buffer:	0x1cf42e0

delete_cb:	Alfa Romeo	164	1993
------------	------------	-----	------

LIST_DELETE SUCCESSFUL

list:	total:	24
	count:	20
	o_size:	36
	buffer:	0x1cf42e0

01:	Audi	Q7	2007
02:	BMW	Z4	2012
03:	BMW	Z4	2008
04:	Buick	Skylark	1997
05:	Chevrolet	Silverado 1500	2003
06:	Dodge	Ram 3500	2002
07:	Ford	Aspire	1994
08:	Ford	Explorer	2007
09:	Hyundai	Equus	2012
10:	Hyundai	Tiburon	2001
11:	Lexus	LS	1994
12:	Lexus	RX	2010
13:	Lexus	SC	2007

14:	Mercury	Capri	1993
15:	Mitsubishi	Truck	1991
16:	Oldsmobile	Bravada	2002
17:	Pontiac	Grand Prix	1968
18:	Pontiac	Grand Prix	1987
19:	Subaru	Legacy	2007
20:	Volkswagen	Cabriolet	1999

delete_cb: Volkswagen Cabriolet 1999

LIST_DELETE SUCCESSFUL

list: total: 24
 count: 19
 o_size: 36
 buffer: 0x1cf42e0

01:	Audi	Q7	2007
02:	BMW	Z4	2012
03:	BMW	Z4	2008
04:	Buick	Skylark	1997
05:	Chevrolet	Silverado 1500	2003
06:	Dodge	Ram 3500	2002
07:	Ford	Aspire	1994
08:	Ford	Explorer	2007
09:	Hyundai	Equus	2012
10:	Hyundai	Tiburon	2001
11:	Lexus	LS	1994
12:	Lexus	RX	2010
13:	Lexus	SC	2007
14:	Mercury	Capri	1993
15:	Mitsubishi	Truck	1991
16:	Oldsmobile	Bravada	2002
17:	Pontiac	Grand Prix	1968
18:	Pontiac	Grand Prix	1987
19:	Subaru	Legacy	2007

LIST_DELETE SUCCESSFUL (NO CALLBACK)

list: total: 24
 count: 18
 o_size: 36
 buffer: 0x1cf42e0

01:	Audi	Q7	2007
02:	BMW	Z4	2012
03:	BMW	Z4	2008
04:	Buick	Skylark	1997
05:	Chevrolet	Silverado 1500	2003
06:	Dodge	Ram 3500	2002
07:	Ford	Aspire	1994
08:	Ford	Explorer	2007
09:	Hyundai	Equus	2012
10:	Hyundai	Tiburon	2001
11:	Lexus	LS	1994
12:	Lexus	RX	2010
13:	Lexus	SC	2007
14:	Mercury	Capri	1993
15:	Mitsubishi	Truck	1991
16:	Pontiac	Grand Prix	1968
17:	Pontiac	Grand Prix	1987
18:	Subaru	Legacy	2007

LIST_TERM SUCCESSFUL

```
list:      total:      0
          count:      0
          o_size:     0
          buffer:    (nil)
```

The members of the `list` structure are printed out after every operation to indicate success. After `list_init` completes successfully the `list` structure print out indicates that: the total number of objects the list can hold is 24; the number of objects in the list is 0; the size of each object is 36-bytes; and the address of the buffer where the objects will be stored—`0x233c2e0`. (NOTE: the address you see on your system will be different). Next, the contents of 24 vehicle structures are added to the list via a for-loop. Then the contents of the `list` structure is displayed with the count indicating that all 24 vehicle structures were indeed added to the list buffer. The list is now full. Next, the contents of the vehicle structures in the list buffer are displayed with the help the iteration functions: `list_begin` and `list_next`. After that, all the GMC vehicles are targeted for deletion from the list. (Nothing against GMC vehicles). The contents of the list structure are displayed with the count now 21 instead of 24. All three GMC vehicles were deleted successfully from the list as the following listing indicates. Next, two attempts are made to delete vehicles that not in the list. Both, attempts fail as expected. Next we successfully delete the first vehicle (Alfa Romeo) and the last vehicle (Volkswagen) from the list. The importance of these two deletions is to test corner cases: one at the beginning of the list; and the other at the end. The last test involves deleting a vehicle (Oldsmobile) without providing a callback function. The deletion is successful as the last listing of vehicles indicates. However, there is no printout of the vehicle being deleted by the callback function. Finally, the `list_term` and `list_free` functions are called and the test program terminates. Take time to look

at the source code for test program `list_delete` in files `main.h` and `main.c` in folder `util/list_delete/`.

I hope you found this tutorial helpful. If you want a challenge see Appendix A for an exercise.

Appendix A

Exercises:

1. We have added a most of the bells and whistles you would expect for a glorified array. However, there is one I can think of that we left undone. The iterator functions `list_end` and `list_prev` to iterate from the end of the list to the beginning. And a test program for these functions.

Appendix B

```
;-----  
; C/C++ definition:  
;  
; void * memmove64 (void *dst, void const *src, size_t size)  
;  
; passed in:  
;  
; rdi = dst  
; rsi = src  
; rdx = size  
;  
; returned:  
;  
; rax = dst  
;  
; WARNING: this routine does not handle the overlapping source-  
;          destination senario.  
;-----  
;  
global memmove64:function  
memmove64:  
    push    rdi                ; push destination address on stack  
; quadword count = size / QW_SIZE  
    mov     rax, rdx           ; copy value of size parameter to rax  
    xor     rdx, rdx           ; zero out rdx  
    mov     r11, QW_SIZE  
    div     r11                ; rdx:rax = (byte count):(quadword count)  
    mov     rcx, rax           ; copy quadword count to rcx  
    cld  
    ; increment index registers rsi and rdi  
    rep movsq                 ; repeat quadword move operation  
    mov     rcx, rdx           ; copy byte count to rcx  
    rep movsb                 ; repeat byte move operation  
    pop     rax                ; pop destination address off stack  
    ret
```

Appendix C

```
;-----  
; C declaration:  
;  
; int list_init (list_t *list, size_t const o_size);  
;  
; param:  
;  
; rdi = list  
; rsi = o_size  
;  
; return:  
;  
; eax = 0 (success) | -1 (failure)  
;  
; stack:  
;  
; [rbp - 8] = rdi (list)  
;-----  
;  
global list_init:function  
list_init:  
; prologue  
    push    rbp  
    mov     rbp, rsp  
    sub     rsp, 8  
; store rdi (list) on stack  
    mov     QWORD [rbp - 8], rdi  
; list->o_size = o_size  
    mov     QWORD [rdi + list.o_size], rsi  
; buffer_size = o_size * LIST_COUNT  
    mov     rax, rsi  
    mov     rcx, QWORD LIST_COUNT  
    mul     rcx  
; buffer_size = (buffer_size + ALIGN_WITH) & ALIGN_MASK  
    add     rax, QWORD ALIGN_WITH  
    and     rax, QWORD ALIGN_MASK  
; if ((list->buffer = calloc(1, buffer_size)) == NULL) return -1  
    mov     rdi, 1  
    mov     rsi, rax  
    call    calloc wrt ..plt  
    mov     rdi, QWORD [rbp - 8]  
    mov     QWORD [rdi + list.buffer], rax  
    test    rax, rax  
    jnz     .continue  
    mov     eax, -1  
    jmp     .epilogue  
.continue:  
; list->total = LIST_COUNT  
    mov     rax, LIST_COUNT
```

Appendix C (cont.)

```
        mov     QWORD [rdi + list.total], rax
; list->count = 0
        xor     rax, rax
        mov     QWORD [rdi + list.count], rax
; return 0
;     xor     eax, eax
.epilogue:
        mov     rsp, rbp
        pop     rbp
        ret
```

Appendix D

```
;-----  
; C definition:  
;  
; void list_term (list_t *list)  
;  
; param:  
;  
; rdi = list  
;  
; stack:  
;  
; [rbp - 8] = rdi (list}  
;-----  
;  
    global list_term:function  
list_term:  
; prologue  
    push    rbp  
    mov     rbp, rsp  
    sub     rsp, 8  
; store rdi (list) on stack  
    mov     QWORD [rbp - 8], rdi  
; free list buffer memory on heap  
    mov     rdi, QWORD [rdi + list.buffer]  
    call    free wrt ..plt  
; zero out list structure  
    mov     rdi, QWORD [rbp - 8]  
    mov     rsi, QWORD listSize  
    call    bzero wrt ..plt  
; epilogue  
    mov     rsp, rbp  
    pop     rbp  
    ret
```

Appendix E

```
;-----  
; C definition:  
;  
; void * list_add (list_t *list, void const *object);  
;  
; param:  
;  
; rdi = list  
; rsi = object  
;  
; return:  
;  
; rax = address of object in list | NULL  
;  
; stack:  
;  
; [rbp - 8] = rdi (list)  
; [rbp - 16] = (void *addr) = address of object in list buffer  
;-----  
;  
    global list_add:function  
list_add:  
; prologue  
    push    rbp  
    mov     rbp, rsp  
    sub     rsp, 16  
; store rdi (list) on stack  
    mov     QWORD [rbp - 8], rdi  
; if (list->count >= list->total) return NULL  
    xor     rax, rax  
    mov     rcx, QWORD [rdi + list.count]  
    cmp     rcx, QWORD [rdi + list.total]  
    jae     .epilogue  
; void *slot = &list->buffer[(list->count * list->o_size)]  
    mov     rax, QWORD [rdi + list.count]  
    mul     QWORD [rdi + list.o_size]  
    add     rax, QWORD [rdi + list.buffer]  
    mov     QWORD [rbp - 16], rax  
; (void) memmove64(slot, object, list->o_size)  
    mov     rdx, QWORD [rdi + list.o_size]  
    mov     rdi, rax  
    call    memmove64 wrt ..plt  
; list->count += 1  
    mov     rdi, QWORD [rbp - 8]  
    mov     rax, QWORD [rdi + list.count]  
    inc     rax  
    mov     QWORD [rdi + list.count], rax  
; return addr  
    mov     rax, QWORD [rbp - 16]  
.epilogue:  
    mov     rsp, rbp
```

Appendix E (cont.)

pop	rbp
ret	

Appendix F

```
;-----  
; C definition:  
;  
;   size_t list_count (list_t *list);  
;  
; param:  
;  
;   rdi = list  
;  
; return:  
;  
;   rax = list->count  
;-----  
;  
    global list_count:function  
list_count:  
    mov     rax, QWORD [rdi + list.count]  
    ret
```


Appendix G

```
;-----  
; C definition:  
;  
; void * list_begin (list_t *list);  
;  
; param:  
;  
; rdi = list  
;  
; return:  
;  
; rax = list->buffer | NULL  
;-----  
;  
    global list_begin:function  
list_begin:  
; if (list->count == 0) return NULL  
    mov     rax, QWORD [rdi + list.count]  
    test    rax, rax  
    jz      .epilogue  
; list->index = 0L  
    xor     rax, rax  
    mov     QWORD [rdi + list.index], rax  
; return list->buffer  
    mov     rax, QWORD [rdi + list.buffer]  
.epilogue:  
    ret
```

Appendix H

```
;-----  
; C definition:  
;  
; void * list_next (list_t *list);  
;  
; param:  
;  
; rdi = list  
;  
; return:  
;  
; rax = (list->buffer + ((list->index + 1) * list->o_size)) | NULL  
;-----  
;  
    global list_next:function  
list_next:  
; if (list->index >= (list->count - 1)) return NULL  
    xor     rax, rax  
    mov     rcx, QWORD [rdi + list.count]  
    dec     rcx  
    mov     rdx, QWORD [rdi + list.index]  
    cmp     rdx, rcx  
    jae     .return  
; list->index += 1  
    inc     rdx  
    mov     QWORD [rdi + list.index], rdx  
; return (list->buffer + (list->index * list->o_size))  
    mov     rax, rdx  
    mul     QWORD [rdi + list.o_size]  
    add     rax, QWORD [rdi + list.buffer]  
.return:  
    ret
```

Appendix I

```
;-----  
; C definition:  
;  
; void list_sort (list_t *list, int (*sort_cb) (void const *, void const *));  
;  
; param:  
;  
; rdi = list  
; rsi = sort_cb  
;-----  
;  
    global list_sort:function  
list_sort:  
    push    r15  
    mov     rcx, rsi  
    mov     rdx, QWORD [rdi + list.o_size]  
    mov     rsi, QWORD [rdi + list.count]  
    mov     rdi, QWORD [rdi + list.buffer]  
    ALIGN_STACK_AND_CALL r15, qsort, wrt, ..plt  
    pop     r15  
    ret
```

Appendix J

```
;-----  
; C declaration:  
;  
; void * list_find (list_t const *list, void const *key,  
; int (*find_cb) (void const *, void const *));  
;  
; param:  
;  
; rdi = list  
; rsi = key  
; rdx = find_cb  
;  
; return:  
;  
; eax = iter (address of matching object) | NULL  
;-----  
;  
    global list_find:function  
list_find:  
    push    r12  
    push    rsi  
    mov     r8, rdx  
    mov     rcx, QWORD [rdi + list.o_size]  
    mov     rdx, QWORD [rdi + list.count]  
    mov     rsi, QWORD [rdi + list.buffer]  
    pop     rdi  
    ALIGN_STACK_AND_CALL r12, bsearch, wrt, ..plt  
    pop     r12  
    ret
```

Appendix K

```
-----
; C definition:
;
; int list_delete (list_t *list, void const *key,
; int (*find_cb) (void const *, void const *),
; void (*delete_cb) (void const *));
;
; param:
;
; rdi = list
; rsi = key
; rdx = find_cb
; rcx = delete_cb
;
; return:
;
; 0 (success) | -1 (failure)
;
; stack:
;
; QWORD [rbp - 8] = rdi (list)
; QWORD [rbp - 16] = rcx (delete_cb)
; QWORD [rbp - 24] = (void *target)
; QWORD [rbp - 32] = (void *blk_tail)
;-----
;
; global list_delete:function
list_delete:
; prologue
    push    rbp
    mov     rbp, rsp
    sub     rsp, 32
; store rdi (list) and rcx (delete_cb) on stack
    mov     QWORD [rbp - 8], rdi
    mov     QWORD [rbp - 16], rcx
; if ((target = list_find(list, key, find_cb)) == NULL) return -1
    call    list_find
    mov     QWORD [rbp - 24], rax
    test    rax, rax
    jnz     .target_found
    mov     eax, -1
    jmp     .epilogue
.target_found:
; delete_cb(target)
    mov     rcx, QWORD [rbp - 16]
    test    rcx, rcx
```

Appendix K (cont.)

```
        jz         .no_delete_cb
        mov        rdi, rax
        call       rcx
.no_delete_cb:
; void *blk_tail = list->buffer + (list->count * list->o_size)
        mov        rdi, QWORD [rbp - 8]
        mov        rax, [rdi + list.count]
        mul        QWORD [rdi + list.o_size]
        add        rax, QWORD [rdi + list.buffer]
        mov        QWORD [rbp - 32], rax
; if (target == (blk_tail - list->o_size)) goto .dec_count
        sub        rax, QWORD [rdi + list.o_size]
        cmp        QWORD [rbp - 24], rax
        je         .dec_count
; void *blk_head = target + list->o_size
        mov        rcx, QWORD [rbp - 24]
        add        rcx, QWORD [rdi + list.o_size]
; size_t blk_size = blk_tail - blk_head
        mov        rax, QWORD [rbp - 32]
        sub        rax, rcx
; memmove64(target, blk_head, blk_size)
        mov        rdi, QWORD [rbp - 24]
        mov        rsi, rcx
        mov        rdx, rax
        call       memmove64 wrt ..plt
.dec_count: ; list->count -= 1
        mov        rdi, QWORD [rbp - 8]
        mov        rax, QWORD [rdi + list.count]
        dec        rax
        mov        QWORD [rdi + list.count], rax
; return 0
        xor        eax, eax
.epilogue:
        mov        rsp, rbp
        pop        rbp
        ret
```